

# XSEDE

Extreme Science and Engineering  
Discovery Environment

## High performance computing in Python


*Presented by:*

- Steve Lantz - [steve.lantz@cornell.edu](mailto:steve.lantz@cornell.edu)
- Roberto Camacho - [rcamachobarranco@utep.edu](mailto:rcamachobarranco@utep.edu)




# Another peek at how microprocessors work

## Multi-core processors





# Another peek at how microprocessors work

## Multithreading



# Another peek at how microprocessors work

## Single thread

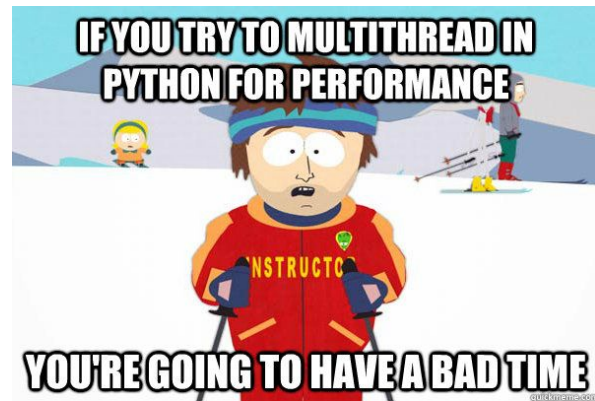


Multiple threads

# Great! How do we exploit this in Python?

Global Interpreter Lock in Python → no multi-threading

- Alternatives:
  - Python multiprocessing module
  - Multi-processing via the OS
  - External application that calls your Python code (e.g., Spark or Hadoop)
  - Code that your Python code calls (e.g. you could have your Python code call a C function that does the expensive multi-threaded stuff).



# Exploiting parallelism in Monte Carlo applications

- Remember, we are trying to model an unknown distribution!
- Works best with a HUGE number of random samples
- Luckily, it's "embarrassingly" parallel
  - Every random sample is independent
- Do the work faster with  $N$  parallel tasks
  - Each does  $1/N$  of the total samples
  - Each must use a different seed
  - Merge results at the end



# Exploiting parallelism in Monte Carlo applications

